

# Bias, variance and regularization

Jan Chorowski  
Instytut Informatyki  
Wydział Matematyki i Informatyki Uniwersytet  
Wrocławski  
2018

# Generalization, the goal of learning

- Problem:
  - We care about the performance on **all the data**
  - We have only a **training sample**
  - **Over-fitting:**  
Too powerful classifier will perfectly interpolate the training data (even the noise in it!) and do poorly on **unseen samples**
  - **Under-fitting:**  
Too weak classifier cannot express the relation in the data, even on **training samples**
- Questions:
  - How to estimate the **generalization** (performance on all data)? -> Honest estimates
  - How to control the capacity of a model?
  - Can we provably ensure good generalization performance -> Learning Theory

# Honest estimates: Hold-out set

Large data case!!!

- Split the training data into two parts:
- Train only on training, then test on testing.
- Often we do a three-way split:
- Then:
  - Train many models on training (different algos, parameters)
  - Use validation to choose best model
  - Test on testing

# Cross-validation

Small data case!!

- Hold-out set makes inefficient data use
- Idea:
  - Divide the data into  $k$  sets ( $\sim 5, 10$ )

For  $i=1..k$

Train on all but the  $i$ -th set

may further split to choose the model...

Test on the  $i$ -th set

Finally:

take the answers on the testing sets and use them to compute the performance measures

- Extreme case: leave-one-out (jackknife) – always use all but one sample to train!

# Bootstrap

- Small data case!!
- Sample with replacement  $m$  samples
  - About 37% will not be selected
- Train on the selected samples
- Test on the remaining ones
- Optionally repeat.

# Bias-Variance: two sources of error!

- The **bias** captures how well our family of functions (hypothesis space) matches the data.
- The **variance** captures how the results of training vary with different samples from the training data

# How to lower the bias?

- Choose more powerful/better models:
  - Understand the data and choose a matching model
  - Describe the data with more attributes
  - **More hidden neurons**
  - Better data transformation
- This usually increases the Hypothesis space

# How to lower variance?

- Get more data (or generate synthetic, e.g. rotate and shear pictures)
- Select only the most important inputs
- **Constrain the models:**
  - Simpler models
  - Regularize the models:  
Assign a probability distribution to the models and choose the most probable ones
- **Average the models**
  - Very powerful
  - Also called “ensemble learning”, boosting, bagging
  - Requires that the models make uncorrelated errors
    - You can even INJECT randomness to decrease the correlation, e.g. *Random Forests*



# Model regularization

- The intuitions:
  - Start with many weights (larger nets train easier!)
  - Choose only the ones that we need. How?
  - Force all the weights to decrease
  - Hope that the necessary ones will remain
  - Subtract a little bit in each training iteration:  
$$\Theta \leftarrow \Theta - \alpha(\nabla_{\Theta} (J) + \beta\Theta) \quad (\text{weight decay})$$
  - Note: this minimizes  $J(\Theta) + \frac{\beta}{2} \sum_j (\Theta_j)^2$
  - Note: usually you don't decay the biases

# Other ideas for NNet regularization

- Choose proper architecture – add or remove neurons or whole layers
- Choose weight decay constants
  - Can also use 1<sup>st</sup> norm, i.e.  $\sum |\Theta_j|$   
Hint: to gradient train approx  $|x| = \sqrt{x^2 + \epsilon}$ ,  
 $\epsilon \approx 10^{-4}$
- Share weights between neurons
  - Example: convolutional networks
- Early stop training
  - Monitor validation error as training progresses. Stop when it starts to increase
- Use dropout -> randomly remove some neurons
- Use weight noise

# Early stopping

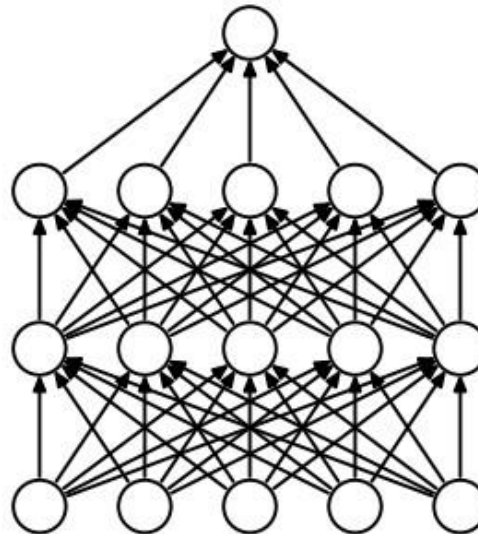
- The net starts with small weights (we initialize it like that)
- Thus it is somewhat linear (all nonlinearities are in the linear range)
- As training progresses the net specializes
- At some point, it over-specializes
- Look for that moment, by monitoring a validation error!

Err rate

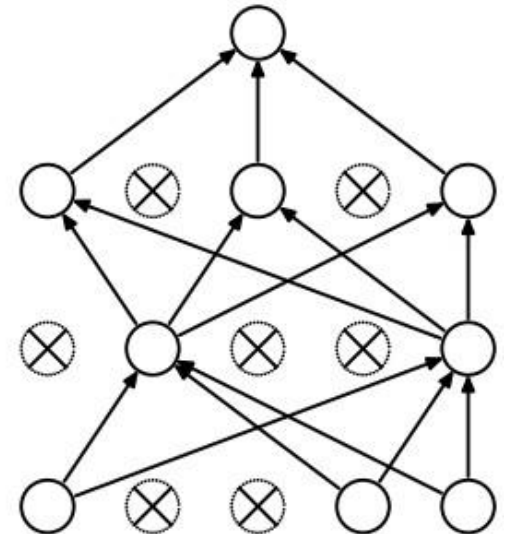
Training iterations

# Dropout

- For each example, select with probability  $p$  which neurons will be used (not dropped out).
- Multiply the outputs of other neurons by  $1/p$  to compensate
- Enjoy!
- Interpretation:
  - Prevents co-adaptation of neurons, as it is harder for neurons to cooperate if any can be dropped-out
  - Trains infinitely many networks, each sharing selected neurons with the other ones



(a) Standard Neural Net



(b) After applying dropout.

# Probabilistic interpretation of Weight Decay

- The gradient step:

$$\Theta := \Theta - \alpha(\nabla_{\Theta} (J) - \beta\Theta)$$

Corresponds to minimizing:

$$J(\Theta) + \frac{\beta}{2} \sum_j (\Theta_j)^2$$

Now try to find a probabilistic interpretation!

# Bayesian approach

1. Make some models more probable than others
2. Set a **prior** probability distribution over  $\Theta$
3. For example:
  - weights are normally distributed  $p(\Theta_i) \sim \mathcal{N}(0, \sigma_\Theta)$
4. Previously we have assumed:  
 $P(Y|X; \Theta)$  i.e.  $y$  depends on  $x$ , with a fixed, but unknown  $\Theta$
5. Now we will treat  $\Theta$  as a random variable too  
 $P(Y|X, \Theta)$  i.e.  $y$  depends on  $x$  and  $\Theta$  which is randomly sampled too

# Bayes theorem

$$p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$

Interpretation: how our estimate of  $A$  changes after seeing  $B$ .

Why?

$$p(A, B) = p(A|B)p(B) = p(B|A)p(A)$$

Then divide by  $p(B)$

# Bayesian approach to ML

- What is the model probability after seeing the examples in set  $S$ ?

$$p(\Theta|S) = \frac{p(S|\Theta)p(\Theta)}{p(S)}$$

How to make predictions? Integrate over all models:

$$p(y|x, S) = \int_{\Theta} p(y|x, \Theta)p(\Theta|S)d\Theta$$

Then

$$E[y|x, S] = \int_y yp(y|x, S)dy$$

But computing  $p(y|x, S)$  is often intractable :(



# Maximum-a-posteriori

- Instead of predicting integrating over all  $\Theta$
- Use the maximally probable  $\Theta$ :

$$\begin{aligned}\Theta_{MAP} &= \arg \max_{\Theta} p(\Theta|S) \\ &= \arg \max_{\Theta} \left( \prod_{i=1}^m p(y^{(i)}|x^{(i)}, \Theta) \right) p(\Theta)\end{aligned}$$

- It's like Max. Likelihood with the extra term.

# Gaussian model MAP

$$\arg \max_{\Theta} \prod_{i=1}^m p(y^{(i)} | x^{(i)}, \Theta) p(\Theta) =$$
$$\arg \max_{\Theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \Theta) + \log(p(\Theta))$$

Now if  $\Theta_j$  are Gaussian with zero-mean:

$$p(\Theta_j) = \frac{1}{\sigma_{\Theta} \sqrt{2\pi}} e^{-\frac{\Theta_j^2}{2\sigma_{\Theta}^2}}$$

Then we recover the weight decay term:

$$-\log p(\Theta_j) \propto \Theta_j^2$$

# Weight decay interpretation

- Weight decay corresponds to adding the weights' sum of squares to the optimization function.
- It can be interpreted as MAP criterion with a prior assumption of a Gaussian weight distribution!
- Other penalties:
  - Sum of absolute values (norm 1) (Lasso penalty), makes weights sparse (many are exactly 0)
  - Mixture of norm 1 and norm 2 (elastic net penalty)

# L2 vs L1 weight regularization intuitions

We can apply two kinds of penalty terms to weights:

- L2 (sum of squares) makes all weights small
- L1 (sum of absolute values) makes weights sparse, i.e. some weights exactly zero, and other larger